The Write & Improve Corpus 2024:
Error-annotated and CEFR-labelled essays by learners of English

Diane Nicholls[♦] **Andrew Caines**[♣] **Paula Buttery**[♣]
[♦] Cambridge University Press & Assessment, U.K.
[♣] ALTA Institute, Computer Laboratory, University of Cambridge, U.K.
[♦] `diane.nicholls@cambridge.org`
[♣] `firstname.secondname@cl.cam.ac.uk`

Abstract

We present a new annotated corpus of written learner English derived from essays submitted to the learning platform Write & Improve (W&I). Users of W&I are presented with automated scoring and feedback on grammatical errors, and are encouraged to act on their error feedback, submitting multiple versions of their essays for any given prompt. We build the corpus on this interplay between *users* and *prompts*, collecting sets of essays submitted by users for a selected list of 50 popular prompts. The prompts include 20 aimed at beginner learners of English, 20 aimed at intermediate learners, and 10 at advanced learners. This distribution reflects the greater use of W&I by beginner and intermediate learners of English. We ensured that the prompts were not likely to elicit personal information and covered a broad range of tasks and topics. This list of prompts enabled us to identify 5050 essay sets written by 766 users, forming the basis for the Write & Improve Corpus, which is being made available on the ELiT website for non-commercial use. We describe the steps we took to ensure the corpus contains appropriate texts, does not include personal information, and comes with annotations relating to Common European Framework of Reference (CEFR) levels and grammatical errors. All essays were submitted between 2020 and 2022 by registered users of W&I who have supplied their first language (L1) in an optional questionnaire. In total, there are more than 23K essays containing more than 3.5 million word tokens. The final versions of each essay set amount to 762K word tokens. There are 22 different L1s in the corpus, with the most common being Spanish, Portuguese, Japanese, Arabic and Vietnamese. We present some descriptive statistics for the corpus, and consider some research use cases.

1 Introduction

Write & Improve (W&I) is an educational technology platform aimed at learners of English, developed by English Language iTutoring (ELiT) – a subsidiary of Cambridge University Press & Assessment (CUP&A). It offers automated grading and error feedback on submitted texts, maintains user profiles, and operates premium subscription options, +Class View and +Test Zone. It has hundreds of thousands of registered users and a global reach, with more than 400K essays being submitted each month, and unique users in Europe, Asia, Africa and the Americas.

Normally, users select from the prompts[1] provided by W&I which are aimed at beginner, intermediate

[1] Key terms –

- USER: a learner using the Write & Improve site.
- PROMPT: a question prompting the composition of an essay in response.
- SUBMISSION or TEXT: a distinct version of an essay written in response to a prompt, deliberately submitted by the user (i.e. not a draft) for auto-marking.
- USER-PROMPT SET: all submissions made by a single user in response to a single prompt. These sets can range from a size of 1 upwards; we only include sets of 2 or more essays in the Corpus.

or advanced level learners of English[2]. The user writes and submits a text of the required length and receives auto-marker assessment in the form of an estimated CEFR level, along with sentence highlighting, grammatical error feedback, a prompt-relevance score, and a progress chart showing the auto-marker scores for the user over their last 10 submissions. The user may then edit and re-submit the text repeatedly, making incremental changes and receiving updated scores and feedback.

We refer to each series of texts submitted by a single user in response to one prompt as a 'user-prompt set'. These sets are the building blocks of this new corpus, THE WRITE & IMPROVE CORPUS 2024, which contains 5050 essay sets written by 766 users with 22 L1s. The dataset contains 23K distinct essays and more than 3.5 million word tokens (counting white-space separated tokens). We describe our motivation and the composition of this new data release below, before providing some research use cases and future plans for the corpus.

Please note that this paper describes the **second version** of the Corpus. The first was released in October 2024 and contained only the 5050 final versions of each essay set, the original and corrected essays in plain-text format with an essay ID so that alignment could be executed at the essay-level. This new second version contains every version from each essay set, along with CEFR grades and other metadata relating to the user and their first language, as well as the prompts and essays in both plain-text and M2 formats.

2 Related work

2.1 Annotated learner corpora

The two existing datasets which are most similar to this new dataset are the FCE Corpus (Yannakoudakis et al., 2011) and the BEA 2019 Shared Task corpus (Bryant et al., 2019). The FCE Corpus contains a sample of 1238 essays from the Cambridge Learner Corpus (CLC), written in the years 2000-1 by candidates for the First Certificate of English (FCE) exam. It is the only section of the CLC for which the full texts are publicly available (cf. the Open CLC available on Sketch Engine which only allows concordance searching), and is a prominent dataset within the educational NLP community, having been used to benchmark various grammatical error detection (GED) and grammatical error correction (GEC) models. It contains expert grammatical error annotations (Nicholls, 2003) and grades, and has recently featured in the GEC Shared Task 2019 (described below) and in a Universal Dependencies annotation project (Berzak et al., 2016).

BEA is the annual workshop on the 'Innovative Use of NLP for Building Educational Applications', organised by SIGEDU, the Association for Computational Linguistics Special Interest Group for Building Educational Applications. In 2019 the BEA Shared Task involved GEC on English texts. The new data released for the task included 3700 W&I texts, which were error-annotated by the ELiT human annotator ('humannotator') team, and amounted to 801K words. In addition, scores from the W&I auto-marker were used to assign each text to a broad CEFR level: A, B or C. In total, 24 teams took part in the shared task, and the results are presented in the overview paper (Bryant et al., 2019) as well as the shared task website.

There are other error-annotated English learner corpora, namely NUCLE, JFLEG and Lang-8, but each has its own characteristics compared to the kind of corpora reviewed above: NUCLE texts are written by a restricted age (university students) and language background (Singaporean students) (Dahlmeier et al., 2013), JFLEG errors were annotated in a different way (as fluency re-writes, rather than the minimal grammar edits applied in most other learner corpora so far) (Napoles et al., 2017). Lang-8 error annotation is user-generated and known to be of variable quality (Mizumoto et al., 2012; Tajiri et al., 2012).

As described by Bryant et al. (2023), English is the best resourced language in terms of learner data, but annotated corpora for other languages do exist. Some featured in the MultiGED-2023 shared task at NLP4CALL on Multilingual Grammatical Error Detection, including GECCC for Czech, Falko-MERLIN for German, MERLIN for Italian, and SweLL for Swedish (Volodina et al., 2023).

The continuing interest in the existing FCE and BEA Shared Task datasets means that it is timely to release a new, high quality, expert-annotated corpus including texts written by learners from a wide range of proficiency levels and different first languages.

[2]Alternatively, users may upload their own prompt and write a response accordingly, also receiving a mark and grammar feedback. In this corpus we only include data written in response to W&I prompts.

2.2 Automated assessment

Automated assessment of essays, or 'automated essay scoring' (AES), is a long-running area of research and one which will be further enabled by this new dataset. The Corpus will enable the training and evaluation of research models for English AES under a multi-level approach in which the essays range from A1 to C2 on the CEFR scale (as opposed to having a set of essays written for an examination targeting a specific CEFR level). This kind of marking is seen as offering helpful indications of progress for learners of English practising their writing, e.g. on Write & Improve.

Previous work on automated marking of learner English essays included feature-based statistical classifiers (Page, 1968; Burstein et al., 2004; Andersen et al., 2013; Yannakoudakis et al., 2018), and subsequently neural network models (Alikaniotis et al., 2016; Riordan et al., 2017; Andersen et al., 2021). In common with the general trend in NLP as a whole, pre-trained large language models (LLMs) are now being applied to the task of essay scoring, with mixed success (Jeon and Lee, 2023; Mizumoto and Eguchi, 2023; Benedetto et al., 2024 (submitted).

This second version of the Corpus features CEFR level annotations for the final versions of each essay set. Please see section 4 for more detail.

2.3 Grammatical error correction

Grammatical error correction (GEC) is another long-standing research task in the field of natural language processing – along with its binary counterpart, grammatical error detection (GED) – and will also be facilitated by this new dataset. Again, the bulk of prior work has been on English, and includes systems trained on large amounts of annotated data in a supervised fashion, as well as more recent approaches involving pre-trained LLMs and 'in-context learning' with zero-shot or few-shot examples.

Machine translation systems pioneered a general purpose approach to GEC, at first with statistical and later neural models (Brockett et al., 2006; Junczys-Dowmunt et al., 2018; Yuan and Bryant, 2021). More recently, edit-based approaches have been proposed in which corrections are applied on a sequence labelling (Omelianchuk et al., 2020) or sequence-to-sequence basis (Stahlberg and Kumar, 2020). Bryant et al. (2023) offers a comprehensive survey of the history and current state of GEC.

Recent studies indicate that LLMs can be prompted to generate corrected versions of ungrammatical inputs. Coyne et al. (2023); Fang et al. (2023); Loem et al. (2023) all perform evaluation of English GEC by OpenAI's LLMs such as ChatGPT, GPT-3.5 and GPT-4. They find that the GPT$*$ models set new state-of-the-art (SOTA) performance on the JFLEG dataset in which the annotators were permitted to carry out naturalistic fluency rewrites (Napoles et al., 2017). However, they perform worse than SOTA on CoNLL-14 and W&I+LOCNESS, which are much larger datasets that were annotated on the basis of minimal edit corrections. Davis et al. (2024) replicated these results with GPT models and in addition open-source models such as Stable Beluga 2, FLAN-T5 and Falcon. Their findings on the wording of the prompt and the number of examples to include in the prompt were mixed, with some models responding better to different prompts and different numbers of examples.

This new corpus will allow further research into GEC for English, enabling further comparison of SOTA systems and LLMs on data annotated by experts. Indeed, the final versions of each essay set have already featured in a shared task on multilingual GEC (MultiGEC-2025) – see section 4.6.

3 Corpus Design & Text Selection

The overarching design principle for this corpus is to collate many texts from a few popular prompts written by as few users as possible. This means that researchers can, for example, study learner progression while holding task and topic relatively controlled – which is desirable since there are known effects of task and topic which interfere with comparisons of learners at different levels (Buttery and Caines, 2012; Caines and Buttery, 2017).

As a starting point, we selected 50 prompts from those which have been most frequently answered by W&I users. To do so we retrieved the 143 active prompts which had been answered at least 5000 times as of March 2021. Table 1 shows how these were distributed across the three W&I difficulty levels – beginner, intermediate, advanced – and how we sampled from these levels, aiming to select 50 prompts

Difficulty level	N. prompts with >=5K responses (Mar 2021)	Prompts to include in Corpus
Just for fun!	4	0
Beginner	52	20
Intermediate	53	20
Advanced	34	10
Total	143	50

Table 1: The Write & Improve prompts with at least 5K submissions as of March 2021, and the sample size of prompts from each difficulty level.

in total. We did not include the 'just for fun' questions[3]. There are fewer submissions to advanced level questions on W&I, and so we only selected 10 of these, as opposed to the 20 beginner and 20 intermediate questions.

In addition to selecting prompts based on popularity, there were other considerations we took into account: how recent they are, whether they add variety in terms of topics covered, and how prototypical they are in terms of task types. We excluded some popular prompts because they tend to elicit personal information – something we avoid including in the corpus.

Our intention was to sample 100 user-prompt sets for each of the 50 selected high-frequency prompts. This would give us a total of 5000 user-prompt sets at the core of our corpus design. The finalised dataset approximately meets this specification, with some variation around the number of user-prompt sets per prompt (see Appendix A for details).

3.1 Data sampling

For each of the 50 selected prompts, we identified user-prompt sets which meet the following criteria:

- The user has supplied their first language (§3.2).
- The user has returned to W&I on at least one subsequent date after first use. This ensures that user engagement with the platform is not just a one-off encounter.
- The user has been scored at least 3 out of 5 by the automated prompt-relevance marking model for at least half of all the essays they have submitted to W&I.
- Essays should be in response to one of the 50 selected high-frequency prompts. The essay(s) submitted by the user for this prompt are referred to as a USER-PROMPT SET.
- The user-prompt set contains at least two versions of an essay for any given prompt. This enables research into the revisions which are made to original essay versions in response to automated error feedback.
- The first version of the user-prompt essay set was submitted no earlier than 1st December 2019, and the final version was submitted no later than 31st December 2022.
- Essays should not be below the minimum word count required by the prompt, and should not be longer than 600 words (white-space delimited).
- Different versions of the essays should not involve wholesale changes (§3.3).
- The essays should not be invalid responses to the given prompts, for any of several possible 'rejection' reasons (§3.4).

[3] The 'just for fun!' questions include, for instance, "If you could be any animal you choose for one day, what animal would you choose? Why?" and "Here in the UK, on 31st December, we celebrate a New Year. In big cities there are fireworks, and around the country families and friends spend time together to welcome the New Year. How do people celebrate the New Year in your country?"

Variable	Value
Prompts	50
Users	766
User-prompt sets	5050
Maximum user-prompt sets per prompt	104
Minimum user-prompt sets per prompt	92
Mean user-prompt sets per prompt	101
Median user-prompt sets per prompt	102
Essays	23,216
First versions	5050
First versions	5050
Intermediate versions	13,116
Total words (white-space delim)	3,551,269
Maximum words per essay	575
Minimum words per essay	22
Mean words per essay	153
Median words per essay	154
Words in first versions	742,319
Words in last versions	762,186
Words in intermediate versions	2,046,764

Table 2: Write & Improve Corpus 2024, descriptive statistics.

With the essays organised into user-prompt sets according to the restrictions listed above, we attempted to retrieve 100 user-prompt sets for each of the 50 prompts, drawing on as few users as possible – in other words preferring users with the most user-prompt sets for our 50 selected prompts. This decision enables us to study user behaviour and variation more closely, as we have a smaller user pool than we might have had if we had selected user-prompt sets randomly. In the event, due to the complexity of meeting all of the above criteria, we do not have 100 user-prompt sets for every prompt: 41 of 50 prompts have 100 or more user-prompt sets. Further statistics are given in Table 2, including the total number of essays collected as a result of the 5050 user-prompts in the Corpus, and the total number of words (counted through white-space delimitation).

3.2 Users' first languages

Registered W&I users can optionally inform us of their first language (L1) in a profile questionnaire. The L1 field used to be a free-form text box but was standardised as a drop-down menu of 120 languages in 2022. In some cases the information may be an over-simplification as, for example, the user may in fact have been bilingual or multilingual from birth or childhood, the user has to choose 'Spanish', 'Portuguese', 'Chinese', *etc*, even if in fact they speak a variety of Mexican Spanish, Brazilian Portuguese or Shanghainese, and their true language could be missing from the list of 120 most spoken languages – hence they might choose "other". But the information, even if it is only indicative of L1, is useful for grouping users into language families and investigating effects and potential influence, or 'transfer', from their first language.

We identify the most common L1s amongst W&I users, and select 22 to work with. They are listed in Table 3 along with the number of users featuring in the Corpus, and the top-level language family given in the Glottolog linguistic database (Hammarström et al., 2024). The distribution of users by first language in the Corpus largely reflects that of W&I as a whole, with L1 Spanish and Portuguese the most common, and Indo-European L1s predominant. As shown in the table, there is a good diversity of L1s from other language families, with seven represented in total. In terms of lower level language families, it is apparent

Language	Family	N.users	Language	Family	N.users
Spanish	Indo-European	65	Dutch	Indo-European	36
Portuguese	Indo-European	60	Italian	Indo-European	36
Polish	Indo-European	57	Hindi	Indo-European	30
Vietnamese	Austroasiatic	57	Persian	Indo-European	19
Japanese	Japonic	53	Burmese	Sino-Tibetan	18
German	Indo-European	52	Czech	Indo-European	18
Arabic	Afro-Asiatic	51	Thai	Tai-Kadai	13
Mandarin Chinese	Sino-Tibetan	49	Indonesian	Austronesian	11
Turkish	Turkic	49	Catalan	Indo-European	2
Romanian	Indo-European	46	Swedish	Indo-European	2
French	Indo-European	41	Ukrainian	Indo-European	1

Table 3: Write & Improve Corpus 2024, users by first language in descending order of frequency, with language family information from Glottolog (Hammarström et al., 2024).

that Romance languages are the most common (Spanish, Portuguese, Romanian, French, Italian, Catalan) with a total of 250 users, followed by Germanic with 90 (German, Dutch, Swedish), and Slavic with 76 (Polish, Czech, Ukrainian).

This range of L1 backgrounds is a welcome addition to the set of English learner data in the public domain: previous corpora either do not feature such information (e.g. the Write & Improve data included in the BEA 2019 shared task (Bryant et al., 2019)), do not have such a range of L1s (e.g. the FCE Corpus (Yannakoudakis et al., 2011) and NUCLE (Dahlmeier et al., 2013)), or have related but not L1-specific information – like EFCamDat which features nationalities only (Geertzen et al., 2014).

3.3 Incremental essay updates

Our primary interest is in process writing – a method by which the author creates a first version of their essay and works to refine it in further iterations. W&I was designed as an aid in this regard, offering grammatical error feedback, highlighting potentially problematic sentences, and estimating a holistic score mapped to the CEFR framework.

We are interested in the way that W&I users revise their essays from one version to the next. We therefore decided to exclude user-prompt essay sets featuring wholesale changes. There is no clear dividing line between process writing and wholesale revision – rather, it is a continuum. We operationalised this continuum using the edit distance between each version of an essay within a user-prompt set. We used optimal string alignment to calculate Damerau–Levenshtein distances, and excluded those with a distance less than 0.35.

This threshold was based on random inspection of W&I data, looking to retain the kind of essay revisions we are interested in (process writing) while excluding those which involve more wholesale changes. We wanted to be as strict as possible whilst retaining as much data as possible to fulfill our other aims – i.e. having sufficient user-prompt sets for each prompt, including only those users for whom we have L1 information, and minimising the number of different users.

3.4 Essay rejection reasons

Essays were reviewed and rejected by ELiT human annotators for any of a number of reasons which made them unsuitable for inclusion in the Corpus. The most common reason was the essay being off-topic for the given prompt. The next most common was the essay containing personal information, which we wanted to avoid releasing in the Corpus, followed by inappropriate content in the essay, and miscellaneous other reasons. In a first round of review, human annotators checked all first and last versions of the user-prompt sets: in other words, 10,000 essays representing 5000 user-prompt sets were checked for appropriateness for inclusion in the Corpus. Of these, 511 essays relating to 307 distinct user-prompt sets were rejected and replaced with new data.

A further 458 user-prompt sets were put forward for inclusion, including both first and last versions of each set: therefore 916 (458*2) essays were reviewed by the same human annotators, and this time 39 essays representing 24 unique user-prompt sets were rejected as inappropriate. The remaining 5127 user-prompt sets were provisionally approved for inclusion in the Corpus (5000 – 307 + 458 – 24 = 5127). The overall rejection rate of 331 user-prompt sets from the 5458 which were inspected is 6.06%.

Subsequently, further user-prompt sets were excluded from the Corpus when the annotators noticed inappropriate use of W&I by two specific users – one for obscenities, the other was evidently not using W&I for language learning purposes. These were prolific users who had a total of 77 user-prompt sets in the provisional Corpus. Therefore we were left with a total of 5050 approved user-prompt sets.

3.5 Final essay selection

This section has been about the design of the Write & Improve Corpus 2024, describing the method of text selection (and rejection). Descriptive statistics for the approved dataset of 5050 user-prompt sets are given in Table 2. 40% of these user-prompt sets were written in response to Beginner level prompts; another 40% were written in response to Intermediate prompts, and the remaining 20% were written in response to Advanced prompts. It is important to note that the level of the prompts may differ from the assessed level of the essays produced by the learners.

In the next section we describe the steps taken to prepare the Corpus for public release, including the types of human annotation applied to the dataset.

4 Data Annotation & Preparation

The first and final essays in each of the 5050 approved user-prompt sets were corrected for grammatical errors by the ELiT annotation team. This means that a total of 10,100 essays have been reviewed for grammatical errors and corrected if any were found (some essays had no errors). The annotation team are all experienced teachers of English as a Foreign Language (EFL). Some are also teacher trainers, EFL materials writers, and examiners.

In addition to error correction, the final versions of each of the 5050 user-prompt sets were marked by the ELiT annotation team, giving an indicative label from A1 to C2+ within the well-established CEFR proficiency framework[4]. We first describe the marking process in §4.1, before discussing the grammatical error annotation in more detail (§4.3).

4.1 Essay grading

The final version of each user-prompt set has been graded by annotators who have been trained and certified as Cambridge examiners but have had no operational experience or onward training. Indeed, the grading exercise is not one that is used in operational settings, since we asked the annotators to grade the essays on a 12-point scale which could then be mapped to CEFR levels with 'plus' gradations: i.e. A1, A1+, B1, B1+, B2, etc. A 13th zero point on the scale was also available – indicating texts which are below A1 level – but it was not used by the annotators. Note also that C2+ was not used by the annotators, so the observed grades in the dataset range from A1 to C2. This means that we have CEFR grades which can be compared to the scoring scale used by the W&I auto-marker, allowing us to measure human-machine agreement and investigate any wide divergences.

Table 4 shows the distribution of grades in terms of CEFR levels. It is clear that there are very few essays at A1 or A1+ level, some at A2, C1+ and C2 levels, with the bulk of the grades falling between A2+ and C1. Note how this unimodal distribution contrasts with the prompt levels responded to by the users – these were two-fifths beginner level, two-fifths intermediate, one-fifth advanced. Instead, beginner level (A*) essays constitute only one-fifth of the 5050 final essays, advanced level essays (C*) make up one-tenth, with the remainder judged to be intermediate (B*) essays (approx. 7-tenths).

[4]CEFR: Common European Framework of Reference for Languages.

CEFR level	Count	Prop (%)	CEFR level	Count	Prop (%)
A1	1	0.0	B2	919	18.2
A1+	11	0.2	B2+	529	10.5
A2	202	4.0	C1	370	7.3
A2+	846	16.8	C1+	102	2.0
B1	959	19.0	C2	38	0.8
B1+	1073	21.2	C2+	0	0.0

Table 4: Write & Improve Corpus 2024, counts and proportions (%) of essays assigned to each CEFR level by human annotators.

Importantly, marking essays in this way – assigning them to a raw CEFR level across the scale without recourse to subscales or other marking criteria – is unorthodox and unlike operational procedures followed at CUP&A. In addition, normally, a learner's CEFR level would be estimated through several datapoints and testing of multiple skills (listening, speaking, reading, writing) rather than from a single essay alone. Nevertheless, the grades give an indication of the user's proficiency and give researchers a way to train and evaluate auto-markers for writing practice platforms, where approximate estimates of proficiency are helpful to learners.

Across the set of 5050 final essay versions, these indicative human grades are within one CEFR level of the the W&I auto-marker grade 92.1% of the time, and beyond two CEFR levels only 0.4% of the time, with a root mean square error of 1.44 (where a difference of 2.0 is equal to one CEFR level). In other words, on average the auto-marker is well within a CEFR level's difference from the human grades. The auto-marker is on the whole more generous than the human markers: awarding a higher grade to an essay in 50.5% of cases, a lower grade in only 19.6% of cases, and the remaining 29.9% involving exact agreement.

The difference in mark distributions is shown in Figure 1: the middle and bottom facets are directly comparable, since they represent the final versions of each user-prompt set only; the top facet represents all essays in the Corpus. It is apparent that (a) the auto-marker in general uses a greater range of CEFR levels, and that (b) the human grades are shifted slightly leftwards compared to the auto-marker – indicating the greater lenience of the auto-marker, on the whole.

By way of comparing how individual essays are marked by humans and the W&I auto-marker, Figure 2 contains scatter plots – one facet for each prompt difficulty level – of the auto-mark (x-axis) against the human mark (y-axis). When the human mark is higher than the auto-mark, the datapoint is coloured purple; when it is lower the datapoint is green, and when it is the same the point is yellow. It is apparent that the leniency of human marking is dynamic and does not apply uniformly across the proficiency scale. For instance, the lowest auto-marks (2 to 4) are more likely to be marked higher by the human examiners; the middle marks (5 to 10) are more likely to be marked lower, hence the denser mass of green dots; and the highest marks (11 to 13) are always marked lower by the humans. There is also an apparent effect of prompt difficulty level, whereby essays responding to the advanced prompts are less likely to be marked more highly by the humans than those written in response to the beginner and intermediate level prompts.

For experimental work in automated essay scoring, we provide the essays in pre-defined train/dev/test splits, using the same assignment as we used for the MultiGEC-2025 shared task (§4.6). Note that the grades for the test set are not yet publicly available. This is to prevent them being incorporated into the web-crawled datasets used to pre-train LLMs – so that the results of any essay scoring experiments with LLMs remain a fair reflection of their capabilities on unseen data. We intend to set up an online evaluation platform so that test set predictions may be submitted and performance measures obtained. We will announce any future development in this respect in relevant channels and in a new version of this paper. In any case, reporting performance on the dev set is encouraged, whether alongside results on the test set or as an interim measure (Dodge et al., 2019).

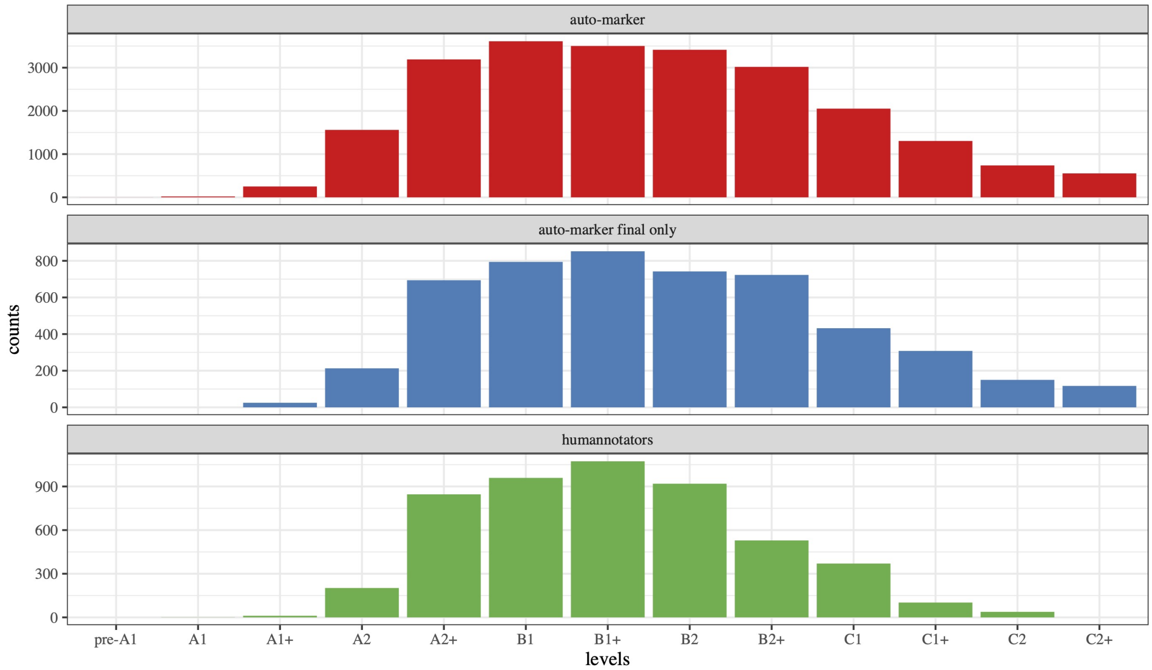

Figure 1: Counts of CEFR levels assigned to essays in the W&I Corpus 2024: the x-axis shows the 13 available grade labels (2 for each of the 6 CEFR levels, plus a 'pre-A1' label), the top facet shows the number of texts labelled at each level according to the auto-marker – for every essay in the Corpus (n=23,216), the middle facet shows the distribution of auto-marks for the subset of 5050 final versions for each user-prompt set in the Corpus, and at the bottom are the human-annotated levels for those 5050 final essay versions.

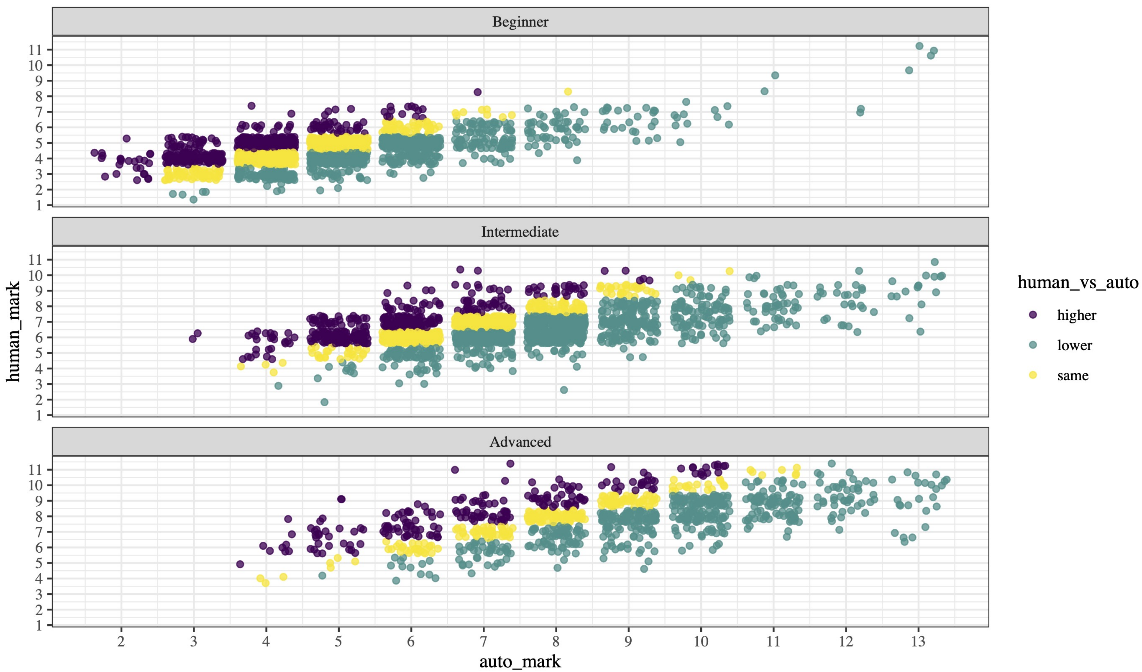

Figure 2: Scatter plots of auto-marks (x) versus human-marks (y), with essays scored more highly by the human markers coloured purple, those scored less highly coloured green, and exact agreement being in yellow. Points are 3/4 opaque and 'jittered' around the vertical to avoid over-plotting.

In Figure 3 we show the difference in auto-marks between the first and final essay versions within each user-prompt set in the Corpus. It is apparent that in the majority of cases, there is an increase in marks between first and final versions of the user's essay for any given prompt. The observation is consistent across prompt difficulty levels, as is the tendency for a small number of user-prompt sets to deteriorate between first and final versions. However, these red points – the ones with a negative difference between first and final version – are skewed towards the sets with lower numbers of essay versions. In other words, the users who persevere more than the norm, with high numbers of essay versions in a user-prompt set, are highly likely to see improvements in their auto-marks by the time of their final essay version. This highlights (a) the benefits of the W&I feedback mechanisms, (b) the importance of reflection and iteration for the learner.

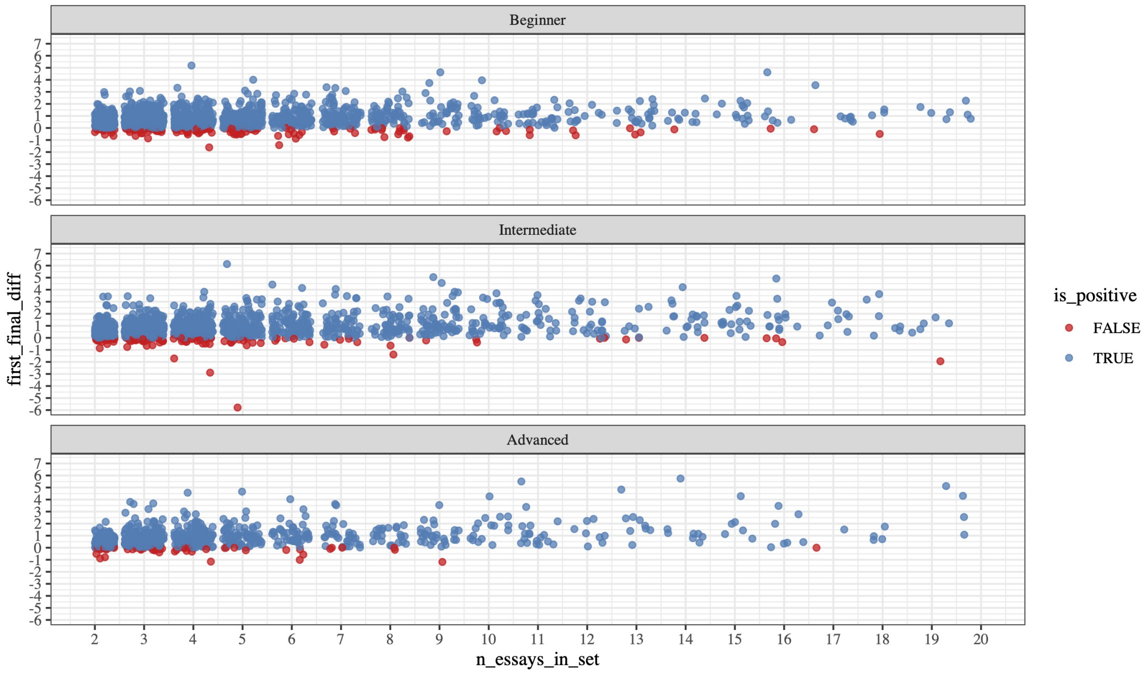

Figure 3: The difference in auto-marks between the first and final essay versions within each user-prompt set in the Write & Improve Corpus 2024 ($markDiff = finalMark - firstMark$). Points coloured blue are those with a positive difference (i.e. an increase) between first and final mark; the red points are those for which there is a negative difference. The numeric difference from -6 to +7 is on the y-axis. The number of essays in the user-prompt set is shown on the x-axis. To avoid over-plotting, the points are 'jittered' randomly around each vertical line rising from an x-axis tick. There is a separate facet for each prompt difficulty level, Beginner, Intermediate, Advanced.

4.2 Learner proficiency level

Note that the true English proficiency levels of each user are not known to us. When they register a W&I account, the user may choose to respond to survey questions about their most recent English exam qualification and the exam they are next targeting – if these questions are applicable – but this still does not certainly inform us about their proficiency level at the time of using W&I, *each time* they use it.

Indeed, even within the Corpus, a user may have essays of different proficiency levels, and each is just an indication of their true level. In reality, a language learner may be assigned a proficiency level based on several data points drawn from tests of different language skills (usually: writing, reading, listening, speaking).

On W&I, users are shown a progress graph of their estimated proficiency level over their most recent essay submissions. Within the Corpus, for instance, we can randomly identify some users with ten essays and plot their auto-marks as shown in Figure 4, which gives a visual impression of the variability of

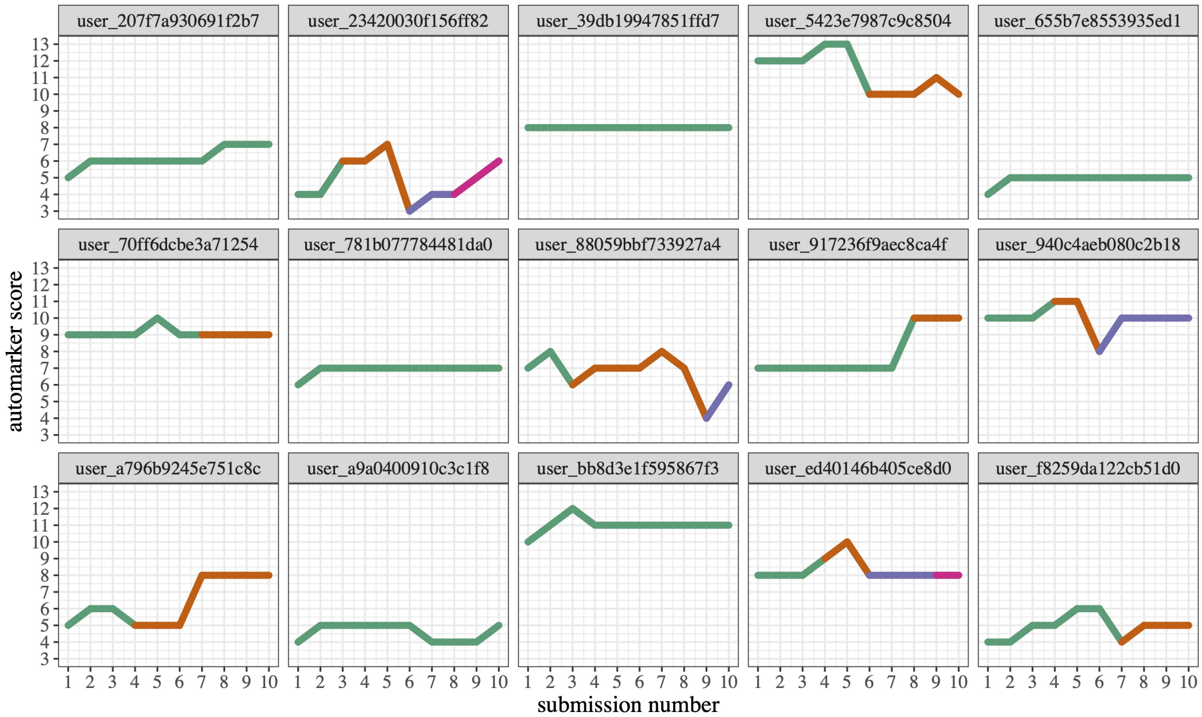

Figure 4: A sample of fifteen users with ten essays in the W&I Corpus 2024, showing their scores from the auto-marker for each essay submission in a separate facet for each user, distinguishing responses to different prompts through changes of colour.

proficiency estimates for a given user based on different texts they produce. Some users score consistently around the same level across their set of ten essays, according to the W&I auto-marker. Others' essays are marked more variably, especially when starting a new essay set for a distinct prompt.

Various, relatively straightforward ways to estimate a user's proficiency level over the set of essays they have submitted to W&I include an overall average, an average based off first versions only, or final versions only, the most recent first versions or the most recent final versions – for instance. More complex, machine learning approaches could be developed based on a user's portfolio of essays submitted to W&I, though we leave this as a matter for future research. In Table 5 we show how the distribution of proficiency levels in the Corpus varies according to a selection of these different methods. Simply taking the most recent auto-marked essay to estimate learner proficiency is liable to misrepresent learner levels due to variability in demands of the task, learner performance, and marking error (this last point is true for both auto-markers and human examiners) – on top of the fact that a single sample of writing is insufficient for an accurate assessment of holistic language ability.

Therefore an average of marks over a larger selection of texts is more likely to yield accurate estimates of proficiency. Taking the average of the marks for all essays submitted by a learner to W&I makes it hard for improvements in proficiency to be rewarded, since the learner's earliest essays will weigh down the overall mean. Focusing on the most recent essays written by the learner would seem the appropriate course of action – whether that be an arbitrary number of recent essays (e.g. the ten most recent, as used here) or essays submitted over a time period such as the last month or quarter. Taking the mean of every essay version would seem the obvious thing to do: alternatively one could calculate the mean of recent first versions only, on the basis that this is a fair reflection of learner proficiency before receiving any feedback from W&I; or, we could calculate the mean based on recent final versions only, as this would indicate the learner's potential proficiency after some feedback to fix basic grammatical errors and improve their writing as much as possible.

In Table 5 we see how these different approaches impact on the distribution of CEFR levels in the

Method	A1	A2	B1	B2	C1	C2
Most recent	1	59	193	260	183	70
Overall mean	2	88	262	241	131	42
First versions	7	134	277	204	117	27
Final versions	0	55	247	259	151	54
Most recent ≤10	2	79	215	271	157	42
Most recent ≤5 first versions	7	116	261	229	126	27
Most recent ≤5 final versions	0	50	219	272	168	57

Table 5: Various methods for estimating learner proficiency levels in the Write & Improve Corpus 2024 for the 766 individuals who feature in the Corpus.

Corpus. Taking the overall mean or any calculation over only first versions moderates the CEFR levels down, whereas the methods involving recent essays or only final versions boost the distribution of CEFR levels upwards. As stated above, a learner's true proficiency level would in reality be estimated based on a number of data points and different skills, rather than simply from their writing, but estimating learner proficiency based on a relatively large sample of writing may still prove to be useful feedback to users of a practice platform such as W&I.

4.3 Grammatical error correction

In common with previous English learner corpora such as the FCE Corpus (Yannakoudakis et al., 2011) and the Write & Improve essays which featured in the BEA 2019 Shared Task (Bryant et al., 2019), the annotators are trained and instructed to carry out grammatical corrections using a principle of minimal intervention – that is, remaining faithful to the learner's perceived intention wherever possible, correcting any errors encountered so that the text becomes adequately grammatical but without rephrasing them unless absolutely necessary. This is a convention known as 'minimal edit' correction, whereas the type of correction involving rephrasing for naturalness and idiomaticity is known as 'fluency' correction (Sakaguchi et al., 2016).

Minimal edits are aimed at facilitating feedback for learning rather than creating a more elegant version of the learner's original text. The pedagogical idea is that it is better to help a learner write correctly what they wrote incorrectly than to help them avoid a language-learning opportunity via circumlocution. For example, **He said me that he would come*, is better corrected to *He said to me that he would come*, than to *He told me that he would come*, because it teaches a generalizable lesson that 'said' requires the preposition 'to', rather than 'teaching' the avoidance of 'said' in favour of 'told'. As a result of this policy, the corrected version of a learner's text may still appear stilted or somewhat unnatural. This is considered fine, as we are error annotating, not editing.

In addition, if the original word or words are incomprehensible even in context, then the annotators had the option to indicate this with an 'interrobang' – meaning that a section of text is problematic but could not be corrected. For instance, compare this original example with its correction:

> ORIGINAL: I realy recommend ti closter, colorfull horbours, beatiful small bays, caves and mountains.
> CORRECTION: I really recommend ‽, colorful harbours, beautiful small bays, caves and mountains.

In this example, edits have been made to fix the misspelling of 'really' and 'colorful'[5]. But the original words, "ti closter", could not be deciphered and therefore no correction has been attempted: instead, an interrobang indicates that this is a problematic section of text.

[5]Note that the annotators are faithful to the original spelling system used: i.e. if American English spellings are in use, then they are not changed to British English.

Rank	Team	GLEU	Precision	Recall	$\mathbf{F}_{0.5}$	Scribendi
1	UAM-CSI	**82.6**	**62.6**	**49.9**	**59.6**	**0.99**
2	Lattice	79.4	54.4	40.2	50.8	0.36
3	baseline	76.4	39.3	42.2	39.8	0.98

Table 6: MultiGEC-2025 results on the Write & Improve Corpus 2024 *development set*: **minimal edits** track.

Rank	Team	GLEU	Precision	Recall	$\mathbf{F}_{0.5}$	Scribendi
1	UAM-CSI	**81.5**	**62.24**	**50.78**	**59.55**	**0.98**
2	Lattice	77.9	58.05	41.6	53.79	0.95
3	baseline	75.15	41.59	41.55	41.58	**0.98**
4	Rum-Cull	60.2	9.63	3.8	7.37	0.34

Table 7: MultiGEC-2025 results on the Write & Improve Corpus 2024 *test set*: **minimal edits** track.

Since there is no declared correction for the issues indicated by interrobangs, for now we reinsert the original tokens into the corrected text, instead of the interrobangs. This means that not attempting to correct those problematic tokens will not be penalised in any evaluation of machine correction systems. In future we may attempt to address this issue with clear identification of the interrobang spans while allowing reference-less evaluation of proposed corrections.

In the first version of the dataset, we released the original and corrected forms of the 5050 final essay versions as plain-text markdown files with an essay ID preceding each essay. This was the format required for the 2025 Shared task on Multilingual Grammatical Error Correction (MultiGEC-2025) held at that year's NLP4CALL Workshop (Masciolini et al., 2025b,a). The essays were randomly assigned into train, dev and test splits at a ratio of 80:10:10. No single user features across splits, and each data split is representative of the overall mixture of L1s and proficiency levels in the main corpus. For more information see the data format section below (§4.6).

The results for MultiGEC-2025 on the Corpus test set are reproduced in Tables 6 and 8 for the development set, and Tables 7 and 9 for the test set (Masciolini et al., 2025a). There were two tracks: one for 'minimal edits' and the other for 'fluency correction'. These represent two different (but overlapping) approaches to GEC in which one can aim to correct the minimum required to make a text grammatical, or one can do that while also making fluency/idiomaticity-related corrections (Sakaguchi et al., 2016). The primary metric used for ranking results in the minimal edits track was $F_{0.5}$; for fluency corrections it was the Scribendi Score (Östling et al., 2024) – in addition GLEU scores were reported (Napoles et al., 2015).

The baseline system for the shared task was based on prompting an LLM with one-shot in-context learning, using the 8B parameter, instruction-tuned version of Llama 3.1 (Llama Team, 2024). UAM-CSI – the winning system in both tracks – was a fine-tuned Gemma 2 model (Staruch, 2025). For reasons stated in Section 4.1 above, the reference corrections for the test set are not being made publicly available. To facilitate GEC evaluation on the test set we intend to set up an online evaluation platform. We will announce any future development in this respect in relevant channels and in a new version of this paper. Note that GLEU scores may currently be obtained from the MultiGEC-2025 CodaLab page (without the need to submit GEC predictions for all 12 languages).

In this second version of the dataset, we also release the same data format for the first versions of each user-prompt essay set – i.e. another 5050 error-corrected essays. This means that in total there are 10,100 error-corrected essays in the Corpus. In addition to that plain-text format, we release the error corrections in conventional M2 format, aligned both at the sentence and essay level. Note that the train/dev/test splits are the same for the first essay versions as the final ones (i.e. the split for the first version is copied from the split for the final version). See section (§4.6) for further details.

By preparing the M2 files, we also obtained error types for the 187,938 edit operations automatically

Rank	Team	GLEU	Precision	Recall	$F_{0.5}$	**Scribendi**
1	UAM-CSI	**82.0**	**63.7**	**47.6**	**59.7**	**0.98**
2	baseline	66.4	21.8	41.0	24.0	0.91

Table 8: MultiGEC-2025 results on the Write & Improve Corpus 2024 *development set*: **fluency corrections** track.

Rank	Team	GLEU	Precision	Recall	$F_{0.5}$	**Scribendi**
1	UAM-CSI	**80.67**	**62.57**	**48.67**	**59.19**	**0.98**
2	baseline	65.86	24.55	40.85	26.68	0.89
3	Rum-Cull	60.2	9.63	3.8	7.37	0.34

Table 9: MultiGEC-2025 results on the Write & Improve Corpus 2024 *test set*: **fluency corrections** track.

Corpus section	**n.essays**	**noop.essays**	**sents/essay**	**noop.sents**	**words/essay**	**edit.rate**
Error-annotated	10,100	0.92%	9.15	30.9%	149	12.5%
First versions	5050	0.03%	9.01	24.5%	147	14.5%
Final versions	5050	0.17%	9.28	37.2%	151	10.6%
auto-A1	194	0.0%	4.65	11.9%	65	27.2%
auto-A2	2285	1.71%	5.51	25.4%	74	18.6%
auto-B1	3281	0.95%	8.06	26.8%	126	15.4%
auto-B2	2670	0.45%	11.0	30.0%	182	11.8%
auto-C1	1256	0.80%	13.3	38.6%	236	8.1%
auto-C2	414	0.24%	15.7	45.7%	307	6.3%
human-A1	12	0.0%	5.58	14.9%	70	27.2%
human-A2	1048	2.19%	5.42	33.3%	67	15.4%
human-B1	2032	1.97%	8.62	30.6%	134	13.9%
human-B2	1448	0.21%	11.6	37.3%	198	9.0%
human-C1	472	1.91%	13.2	55.8%	253	4.5%
human-C2	38	10.5%	13.5	80.0%	285	1.3%

Table 10: The number of essays in different sections of the W&I Corpus 2024, including the whole set of error-annotated essays (n=10.1k), the first and final versions, and the 6 broad CEFR levels as determined by the **auto**-marker and **human** annotators; showing the mean number of sentences and words per essay, the percentage of essays and sentences with no error corrections ("noop" = 'no operation'), and the edit rate per word. Note that the auto-marking section represents all 10.1k error-annotated essays, whereas the human-annotated section relates to the 5050 final versions only.

detected by ERRANT (Bryant et al., 2017) in the 10,100 human-corrected first and final versions of each user-prompt set. This is an overall edit rate of 1 in 12.5 tokens: greater than that reported for the B2 level FCE Corpus; less than the rate reported for an error-annotated corpus of spoken English (see Table 2 in Caines et al. (2020)).

Table 10 shows how the number of words written, the edit rate, and the proportion of 'noop' (no operation = no corrections needed) varies over text version and CEFR level. Firstly it is apparent that the final versions of the essays are on average a few words longer than the first versions, and that grammatical errors reduce – evidenced by the higher noop and lower edit rates in the set of final essay versions compared to first versions. This might be expected through the iterative writing process, but also speaks

Rank	Error type	Count	Proportion
1	R:OTHER	30,295	16.1%
2	R:NOUN	24,733	13.2%
3	R:SPELL	14,924	7.9%
4	R:ORTH	13,261	7.1%
5	M:DET	10,356	5.5%
6	R:PREP	9496	5.1%
7	U:DET	6969	3.7%
8	R:VERB	6734	3.6%
9	R:VERB:FORM	6113	3.3%
10	R:VERB:TENSE	5677	3.0%

Table 11: The 10 most frequent error types in the Write & Improve Corpus 2024 (first & final essay versions), both as counts and proportions of all errors. The initial letter before the colon indicates the edit type (**R**eplacement, **M**issing, **U**nnecessary); the error type is shown after the first colon.

to the language improvements which are supported through the automated feedback features in W&I.

When we look across CEFR levels, the bulk of essays being A2 to C1 levels and the rarity of A1 and C2 essays in the Corpus are apparent – whether judged by the auto-marker or human annotators. On the whole, the noop essay and sentence rates increase as CEFR increases also. The development is not straightforwardly linear, however, and there are signs of U-shaped development, or a plateauing in improvement, around the intermediate points of the CEFR framework. This may well be as learners attempt to write more complex constructions as their own proficiency improves. Again, the number of sentences per essay and the number of words per essay both increase with increasing CEFR level. The edits deemed necessary by the humannotators evidently decrease with increasing CEFR level. The essays judged to be at C2 level by the human graders have the most pronounced reduction in edits and increase in noop sentences and essays.

Table 11 shows the ten most frequent error types in the Corpus overall[6]. The fact that the ‘replacement’ edit type[7] (R) and the ‘other’ error type feature prominently in the list is consistent with previous analysis (e.g. Table 9 (note the caption) and Table 10 in Bryant et al. (2019) and Table 2 in Caines et al. (2020)). The high frequency of other, noun and verb errors indicates the difficulty of the GEC task, since these are some of the error types on which relatively low performance has been reported previously; on the other hand, systems have been relatively proficient in the past at correcting spelling, orthography, determiner and preposition errors (e.g. Tables 10 and 11 in Bryant et al. (2019)).

4.4 Grammatical error detection

For each essay we also release grammatical error detection data, which are automatic error predictions made by the GED rule-based model used in W&I at the time the essays were submitted (in the years 2020-2022).

Note that high precision is prioritised over recall, since in an application setting it is deemed preferable to give accurate error feedback rather than trying to detect every error. This approach is in line with the conventional practice in GED and GEC research to weight precision twice as much as recall when evaluating with the $F_{0.5}$ measure.

In practical terms, this approach is operationalised for W&I based on the errors found in the large Cambridge Learner Corpus (CLC) of learner essays. Error detection and correction rules are derived from n-grams which have a high frequency as errors given all the times they occur. Thus the GED rules

[6] See Table 2 in Bryant et al. (2017) for a full list of error types in the ERRANT taxonomy, with descriptions.

[7] As opposed to the ‘missing’ (M) or ‘unnecessary’ (U) error types for which the appropriate correction is insertion or deletion, respectively.

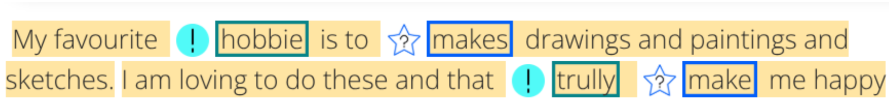

Figure 5: Example of grammatical error detection feedback in the Write & Improve user interface.

token	label
it	c
’s	c
times	i
to	c
goe	i
.	c

Table 12: An example of grammatical error detection data in tabular .conll format: each row is one token; ‘c’ indicates a correct token, ‘i’ indicates an incorrect token. An empty line separates each sentence. Note that there is no header row in .conll files.

are a restricted set which only grows as the set of annotated training data grows: there is no attempt to anticipate errors for which there is as yet no evidence in the CLC. In addition, the GED rule-set used by W&I has been expanded since the essays in the Corpus were written: as a result, error recall is now three times better than it was for the previous version of the system which was in operation between 2020 and 2022 when the essays in the Corpus were written.

Detected errors are highlighted in the W&I interface (e.g. Figure 5), prompting users to reflect on those problematic tokens and correct them in a revised version of their essay. With the previous version of this error feedback system, corrections were not always given to users – the idea being that it is often better for learning if the user reflects and self-corrects. In the new version of the feedback system there is a scaffolding mechanism whereby repeated errors lead to corrections and, where possible, supporting teaching resources, being given.

In the dataset release, GED labels are given on a token-by-token basis in .conll format files – that is, tabular data with one token per row, and errors indicated as binary ‘c’ (correct) or ‘i’ (incorrect) labels, following convention (Volodina et al., 2023). Table 12 shows a made-up example of a sentence in .conll format with GED labels.

In Table 13 we show how well the W&I error detection system performs on the dev and test sets compared to the gold-standard human annotation. It is apparent that (a) precision is much higher than recall – which was an intentional design feature – and that (b) performance is higher for the first over the final versions of the user-prompt essay sets. This indicates that W&I acts as an effective learning tool, enabling users to correct initial errors in the early versions of their essays, and then being less capable of support and error detection for the later versions of their essays – by which time they have been improved and made more linguistically complex, with errors which are in turn more complex and harder to automatically detect with frequency-based detection rules. This supposition is supported by the human annotators’ edit rates being greater for first essay versions over final versions (Table 10).

In future we can add performance statistics for the new version of the W&I error detection system, for which we expect recall to be much higher. Further research is needed to identify the errors which remain out of reach for the detection system, as well as obtaining comparable performance metrics for state-of-the-art neural network GED models. For the dev set we release reference GED data for the first and final user-prompt essay versions in .conll format. We intend to set up an online evaluation platform so that GED predictions for the test set may be submitted and performance measures obtained. We will announce any future development in this respect in relevant channels and in a new version of this paper. We also encourage the reporting of results on the dev set, in line with recommendations for good research practice (Dodge et al., 2019).

Split	Version	Precision	Recall	$\mathbf{F}_{0.5}$
dev	first	0.558	0.145	0.355
dev	final	0.245	0.037	0.116
test	first	0.670	0.139	0.380
test	final	0.319	0.019	0.078

Table 13: Grammatical error detection: precision, recall, $F_{0.5}$ for the W&I feedback system measured against gold-standard annotation.

4.5 Anonymisation, sanitisation & data protection

We took several steps to ensure that personally identifying information does not feature in the Corpus. These were namely: (a) avoiding prompts which are known to elicit personal information; (b) manually reviewing the first and final versions of every user-prompt essay set, checking for personal information among other things, and filtering these essay sets out (§3.4); (c) running an automated check for named entities on all other essays.

We used the spaCy English web Transformer model (v3.7.3)[8] and checked any LOC (location) or PERSON entities found. We did not find any problematic locations. Examples include Cambridge, Central America, Pacific Ocean, Lake Como, Halong Bay, Milky Way, Mars, the Moon.

Most PERSON entities were famous people and therefore were not changed. Examples include Tom & Jerry, Vincent Van Gogh, Napoleon Bonaparte, Clint Eastwood, Harry Potter, James Bond, Leonardo Dicaprio, Harry Styles, Shawn Mendes, Ed Sheeran, Beyoncé, Camilla Cabello, Dua Lipa, Daniel Defoe, Michael J Fox, Emmet Brown, Krzysztof Zalewski (a Polish singer), Bolesław Prus (a movie director), Liam Neeson, Jeremy Irons, Jennifer Lopez, Kayhan Kalhor (a Kurdish Iranian musician), Elon Musk, Cinderella, Santa Claus, Nicole Kidman, Sean Connery, Karl Marx, Ara Malikian (violinist), Anne Frank, Mona Lisa, and Captain America.

If we did not recognise the PERSON entity as a well-known person, we replaced at least one of the names so that there would still be a recognised entity in the text, simply not with the original name. This was a process of pseudonymisation (Volodina et al., 2020): we took care not to change the ethnic origin of the name in case this impacted on system performance, even if preliminary indications are that it does not (Muñoz Sánchez et al., 2024).

We also took steps to sanitise the dataset: to ensure that no offensive terms or rhetoric feature in the Corpus. Firstly, user-prompt essay sets could be rejected for offensiveness by the human reviewers in an early stage of the data preparation process. Some essay sets, and indeed users, were filtered out of the Corpus for this reason (§3.4).

Secondly, we ran a check on all the essays for words in the Surge AI profanity list[9]. This process flagged up several false positives (e.g. essays featuring the words, ‘meat’, ‘loose’ or‘ trash’) but it is best to be cautious. There were two cases to moderate: one involving use of the f-word, which we removed, and another which mentions “*butt* kick runs”. This turns out to be a type of exercise, but appeared in a sentence which mentions "suicide tours" (also about exercise) – nevertheless, we removed this sentence from all versions of the essay, as it could be inadvertently triggering.

Our final data protection step was to delink the identifiers featuring in the Corpus from the application ones generated by W&I. The original W&I identifiers are 36-character alphanumeric strings, with no way to personally identify the user. Even so, as an added layer of data protection and to further distance the public data from private ELiT data, we decided not to release these. Instead, we generated new public identifiers for essays, users and prompts using the `ids` package in R (FitzJohn, 2021). The new IDs are unique 8-character alphanumeric strings prefixed by ‘essay_’, ‘user_’ or ‘prompt_’. As a consequence, the essay, user and prompt IDs carry no significance other than Corpus-internal identifiers.

[8]Reported performance on named entity recognition: precision 0.90, recall 0.90, F1 0.90.

[9]A publicly available word list containing more than 1600 obscene words: `https://github.com/surge-ai/profanity`

4.6 Data format & data splits

In the **first version** of the Corpus, the final versions of each user-prompt essay set were released as plain-text files in the format required by the 2025 Shared task on Multilingual Grammatical Error Correction (MultiGEC-2025) (Masciolini et al., 2025b). The results of this shared task were presented at the 14th Workshop on Natural Language Processing for Computer Assisted Language Learning (NLP4CALL), held at the NoDaLiDa/Baltic-HLT Conference in Tallin, Estonia (Masciolini et al., 2025a).

Twelve languages featured in MultiGEC-2025, and the Write & Improve Corpus 2024 was included as the English language component. The format required for the task is as shown below – parallel files of original and corrected essays, initiated by a metadata line giving the essay ID. The reason for this choice of format is the current prevalance of large language models (LLMs) in GEC research (Caines et al., 2023). Thus we provide the essays according to the way we expect them to be given as input to LLMs – as whole essays without tokenization.

```
### essay_id = 1
Full text of the first essay.

### essay_id = 2
Full text of the second essay.
```

Only the final versions of the user-prompt essay sets were included in this dataset release for the MultiGEC shared task: in other words, the 5050 final essay versions in their original and corrected forms.

To create training, development and test splits from the 5050 essays, we use the R package `rsample` (Frick et al., 2024) The split is 80%:10%:10% and samples are grouped by user, so that no single user's essays feature in more than one data split. There are 4040 essays in the training set, 506 in the development set, and 504 in the test set.

These essays in MultiGEC-2025 markdown format are included in the new release of the Corpus, along with the one-essay-per-line format that we prepared for evaluation purposes, and M2 files for the train and dev splits, output by the ERRANT alignment and error-typing program (Bryant et al., 2017).

In this **second version** of the dataset, we also release the same data format for the first versions of each user-prompt essay set – i.e. another 5050 error-corrected essays. We use the same train/dev/test splits as were established for MultiGEC-2025 by looking up the split for the final essay version, for any given first version, and using that same split for the first version. In total then, there are 10,100 essays in original and corrected form.

The contents of the current version of the Corpus are as follows:

- For the final versions of the 5050 user-prompt essay sets in the Corpus:
  - `.md` markdown, plain-text, non-tokenized essays per MultiGEC-2025 format, original `orig` and corrected `ref1` files for the `train` & `dev` sets, and the original texts but no corrections for the `test` set – note that these five files are equivalent to the first version of the Corpus;
  - `.tmp` one-essay-per-line plain-text essays tokenized with `syntok`, original and corrected files for the `train` & `dev` sets, and the original texts only for the `test` set;
  - `.m2` format file for the `dev` set, showing corrections and edit/error types, obtained from an adapted version of ERRANT (Bryant et al., 2017) as part of the MultiGEC-2025 shared task, and carried out using a UDPipe English model, rather than the spaCy model used conventionally for ERRANT – this was for reasons of consistency in GEC evaluation across the 12 languages which featured in the shared task, as explained in Masciolini et al. (2025a);
  - `local_eval` directory containing instructions and code for local evaluation of GEC hypotheses with ERRANT and GLEU, per the methods used for the MultiGEC-2025 shared task – see the `README` file for further details.
- For the 10,100 error-annotated essays, the first *and* final versions of the 5050 user-prompt essay sets in the Corpus:

  - `.md` markdown, plain-text, non-tokenized essays per MultiGEC-2025 format, original and corrected (`train`, `dev`) and original only (`test`) – in separate files for the `first` and `final` versions;
  - `.tmp` one-essay-per-line plain-text, non-tokenized essays, original and corrected (`train`, `dev`) and original only (`test`);
  - `.m2` format files for the `train` and `dev` sets, showing corrections and edit/error types, both at the sentence and essay level, obtained from version 3.0.0 of ERRANT which in turn uses the small spaCy English model trained on web data – thus note that, though the sentences and essays are assigned to the same train/dev/test splits and appear in the same order, there may be minor differences between these `.m2` files and the ones prepared using the final essay versions for the MultiGEC-2025 shared task (listed above), due to the use of different NLP tools with ERRANT;
  - `.tsv` files which indicate the order of texts in the essay or sentence `.m2` files, for each split of the data, with basic metadata;
  - `.orig` (original) and `.corr` (corrected) parallel plain-text files containing aligned sentences (automatically aligned using Labcoat) for the `train` and `dev` sets, plus original sentences only for the `test` set – this has been one of the conventional formats for GEC in recent years when it operated primarily at the sentence level rather than essay/document level – note that there can be many-to-one, one-to-many and many-to-many sentence alignments depending on the nature of the corrections applied;
  - `.ids` files containing lists of essay IDs for each line in the `.orig` and `.corr` files;
  - `.conll` format files with one token per line and a 'c' (correct) or 'i' (incorrect) label indicating grammaticality for the GED task for the `train` and `dev` sets only – these have been converted from the `.m2` files and therefore tokenisation is carried out by the spaCy model used by ERRANT.

- For all 23,216 essays in the Corpus:
  - `.tsv` with information about the prompts which feature in the Corpus (see Appendix A for a summary);
  - `.tsv` with all texts, untokenized, and all metadata (see Appendix B for more information about the Corpus metadata).

4.7 Use of the Corpus with LLMs & GEC or essay scoring evaluation

We encourage the reader to consider the problem of 'data leakage' with regard to LLMs: whereby NLP datasets are leaked into LLM training datasets via commercial APIs (Balloccu et al., 2024). We therefore ask that the Corpus is only passed to locally-stored LLMs, such as might be downloaded from the Hugging Face Transformers library (Wolf et al., 2020), or to commercial LLMs in such a way that the data is not retained for model training purposes.

It is due to the risk of leakage to LLMs (or indeed, unwanted optimisation by any system) that we do not at present intend to release the reference annotations for GEC or essay scoring for the test set of the Corpus. There is precedent for this from the BEA 2019 Shared Task dataset (Bryant et al., 2019), for which, at the time of writing, the test set references have not been made public. Evaluation of BEA'19 test set data is currently carried out via a Codalab site. Evaluation for the test set of the Write & Improve Corpus 2024 may also be carried out via the MultiGEC-2025 Codalab page, but presently only for the GLEU metric (Napoles et al., 2015). Scripts and instructions for local evaluation with all three metrics used in MultiGEC-2025 (GLEU, $F_{0.5}$ with the ERRANT scorer, the Scribendi Score) are available from a GitHub repository (the scripts for GLEU and ERRANT scoring are also included in the data release; §4.6).

One consequence of the difficulty of working with test set data is that researchers tend to report results on development sets instead or as well. This is a positive side-effect as it is in line with recommended practice for the reporting of experimental results in NLP (Dodge et al., 2019). In addition, the benefit of being able to work with reference files locally (e.g. as is currently the case for the development sets) is the ability to undertake error analysis on system hypotheses. We intend to create a site for online evaluation of GED and essay scoring for the Corpus test set as well; updates about this will follow in due course.

4.8 How to obtain the Corpus

The current second version of the Write & Improve Corpus 2024 may be obtained by visiting the ELiT website, completing the form and agreeing to the licence. Any future versions of the Corpus will be made available on the same website.
Please contact us to stay informed about future updates.

5 Conclusions & Future Work

In this paper we describe the motivation for creating a new corpus of English learner essays annotated with grammatical error corrections and CEFR proficiency labels. The Write & Improve Corpus 2024 contains 23K essays written by 766 users, involving 5050 user-prompt essay sets, and 22 learner L1s. This second version of the Corpus contains various file formats for various purposes; we plan to update the Corpus with further annotation in future. We envisage that the Corpus can support research into grammatical error detection and correction, automated essay scoring, and more besides.

Acknowledgements

This paper reports on work supported by Cambridge University Press & Assessment. We thank Amanda Cowan and the CUP&A Data & Annotation Coordination Committee for their support towards the release of this dataset.

We are grateful to staff at ELiT who have given access to and information about the W&I analytics database, supported text annotation, and helped with data processing: namely (in chronological order), Henry Garner, Scott Thomas, Paweł Stroinski, and David Dursun. We also thank the ELiT annotation team for all their hard work, and Paul Ricketts for his help with final publication.

We thank colleagues in the ALTA Institute and at CUP&A who have given feedback about the proposal. Our thanks to Øistein Andersen for his support and advice; to former group members Mariano Felice and Ted Briscoe for their support of this project, and to Christopher Bryant for discussion and help with the ERRANT toolkit.

References

Dimitrios Alikaniotis, Helen Yannakoudakis, and Marek Rei. 2016. Automatic text scoring using neural networks. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 715–725, Berlin, Germany. Association for Computational Linguistics.

Øistein E. Andersen, Helen Yannakoudakis, Fiona Barker, and Tim Parish. 2013. Developing and testing a self-assessment and tutoring system. In *Proceedings of the Eighth Workshop on Innovative Use of NLP for Building Educational Applications*, pages 32–41, Atlanta, Georgia. Association for Computational Linguistics.

Øistein E Andersen, Zheng Yuan, Rebecca Watson, and Kevin Y. F. Cheung. 2021. Benefits of alternative evaluation methods for automated essay scoring. In *Proceedings of the 14th International Conference on Educational Data Mining (EDM 2021)*, Paris, France.

Simone Balloccu, Patrícia Schmidtová, Mateusz Lango, and Ondrej Dusek. 2024. Leak, cheat, repeat: Data contamination and evaluation malpractices in closed-source LLMs. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 67–93, St. Julian's, Malta. Association for Computational Linguistics.

Luca Benedetto, Gabrielle Gaudeau, Andrew Caines, and Paula Buttery. 2024 (submitted). Assessing how accurately large language models encode and apply the Common European Framework of Reference for Languages. *Computers and Education: Artificial Intelligence*.

Yevgeni Berzak, Jessica Kenney, Carolyn Spadine, Jing Xian Wang, Lucia Lam, Keiko Sophie Mori, Sebastian Garza, and Boris Katz. 2016. Universal Dependencies for learner English. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 737–746, Berlin, Germany. Association for Computational Linguistics.

Chris Brockett, William B. Dolan, and Michael Gamon. 2006. Correcting ESL errors using phrasal SMT techniques. In *Proceedings of the 21st International Conference on Computational Linguistics and 44th Annual Meeting of the Association for Computational Linguistics*, pages 249–256, Sydney, Australia. Association for Computational Linguistics.

Christopher Bryant, Mariano Felice, Øistein E. Andersen, and Ted Briscoe. 2019. The BEA-2019 shared task on grammatical error correction. In *Proceedings of the Fourteenth Workshop on Innovative Use of NLP for Building Educational Applications*, pages 52–75, Florence, Italy. Association for Computational Linguistics.

Christopher Bryant, Mariano Felice, and Ted Briscoe. 2017. Automatic annotation and evaluation of error types for grammatical error correction. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 793–805, Vancouver, Canada. Association for Computational Linguistics.

Christopher Bryant, Zheng Yuan, Muhammad Reza Qorib, Hannan Cao, Hwee Tou Ng, and Ted Briscoe. 2023. Grammatical error correction: A survey of the state of the art. *Computational Linguistics*, pages 643–701.

Jill Burstein, Martin Chodorow, and Claudia Leacock. 2004. Automated essay evaluation: The Criterion online writing service. *AI Magazine*, 25(3):27.

Paula Buttery and Andrew Caines. 2012. Normalising frequency counts to account for 'opportunity of use' in learner corpora. In Yukio Tono, Yuji Kawaguchi, and Makoto Minegishi, editors, *Developmental and Crosslinguistic Perspectives in Learner Corpus Research*. Amsterdam: John Benjamins.

Andrew Caines, Luca Benedetto, Shiva Taslimipoor, Christopher Davis, Yuan Gao, Øistein Andersen, Zheng Yuan, Mark Elliott, Russell Moore, Christopher Bryant, Marek Rei, Helen Yannakoudakis, Andrew Mullooly, Diane Nicholls, and Paula Buttery. 2023. On the application of large language models for language teaching and assessment technology. In *Proceedings of the Empowering Education with LLMs – the Next-Gen Interface and Content Generation Workshop at AIED*.

Andrew Caines, Christian Bentz, Kate Knill, Marek Rei, and Paula Buttery. 2020. Grammatical error detection in transcriptions of spoken English. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 2144–2162, Barcelona, Spain (Online). International Committee on Computational Linguistics.

Andrew Caines and Paula Buttery. 2017. The effect of task and topic on opportunity of use in learner corpora. In Lynne Flowerdew and Vaclav Brezina, editors, *Written and spoken learner corpora and their use in different contexts*. London: Bloomsbury.

Steven Coyne, Keisuke Sakaguchi, Diana Galvan-Sosa, Michael Zock, and Kentaro Inui. 2023. Analyzing the performance of GPT-3.5 and GPT-4 in Grammatical Error Correction. arXiv:2303.14342.

Daniel Dahlmeier, Hwee Tou Ng, and Siew Mei Wu. 2013. Building a large annotated corpus of learner English: The NUS corpus of learner English. In *Proceedings of the Eighth Workshop on Innovative Use of NLP for Building Educational Applications*, pages 22–31, Atlanta, Georgia. Association for Computational Linguistics.

Christopher Davis, Andrew Caines, O Andersen, Shiva Taslimipoor, Helen Yannakoudakis, Zheng Yuan, Christopher Bryant, Marek Rei, and Paula Buttery. 2024. Prompting open-source and commercial language models for grammatical error correction of English learner text. In *Findings of the Association for Computational Linguistics ACL 2024*, pages 11952–11967, Bangkok, Thailand and virtual meeting. Association for Computational Linguistics.

Jesse Dodge, Suchin Gururangan, Dallas Card, Roy Schwartz, and Noah A. Smith. 2019. Show your work: Improved reporting of experimental results. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2185–2194, Hong Kong, China. Association for Computational Linguistics.

Tao Fang, Shu Yang, Kaixin Lan, Derek F. Wong, Jinpeng Hu, Lidia S. Chao, and Yue Zhang. 2023. Is ChatGPT a highly fluent grammatical error correction system? a comprehensive evaluation. arXiv:2304.01746.

Rich FitzJohn. 2021. *ids: Generate Random Identifiers*. Https://github.com/reside-ic/ids, https://reside-ic.github.io/ids/.

Hannah Frick, Fanny Chow, Max Kuhn, Michael Mahoney, Julia Silge, and Hadley Wickham. 2024. *rsample: General Resampling Infrastructure*. R package version 1.2.1, https://github.com/tidymodels/rsample.

Jeroen Geertzen, Theodora Alexopoulou, and Anna Korhonen. 2014. Automatic linguistic annotation of large scale L2 databases: The EF-Cambridge Open Language Database (EFCamDat). In *Selected proceedings of the 2012 Second Language Research Forum*.

Harald Hammarström, Robert Forkel, Martin Haspelmath, and Sebastian Bank. 2024. Glottolog 5.0.

Jaeho Jeon and Seongyong Lee. 2023. Large language models in education: A focus on the complementary relationship between human teachers and ChatGPT. *Education and Information Technologies*, pages 1–20.

Marcin Junczys-Dowmunt, Roman Grundkiewicz, Shubha Guha, and Kenneth Heafield. 2018. Approaching neural grammatical error correction as a low-resource machine translation task. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 595–606, New Orleans, Louisiana. Association for Computational Linguistics.

Llama Team. 2024. The Llama 3 Herd of Models.

Mengsay Loem, Masahiro Kaneko, Sho Takase, and Naoaki Okazaki. 2023. Exploring effectiveness of GPT-3 in grammatical error correction: A study on performance and controllability in prompt-based methods. In *Proceedings of the 18th Workshop on Innovative Use of NLP for Building Educational Applications (BEA 2023)*, pages 205–219, Toronto, Canada. Association for Computational Linguistics.

Arianna Masciolini, Andrew Caines, Orphée De Clercq, Joni Kruijsbergen, Murathan Kurfalı, Ricardo Muñoz Sánchez, Elena Volodina, and Robert Östling. 2025a. MultiGEC-2025 Shared Task at NLP4CALL: Multilingual Grammatical Error Correction. In *Proceedings of the 14th Workshop on NLP for Computer Assisted Language Learning (NLP4CALL)*.

Arianna Masciolini, Andrew Caines, Orphée De Clercq, Joni Kruijsbergen, Murathan Kurfalı, Ricardo Muñoz Sánchez, Elena Volodina, Robert Östling, Kais Allkivi, Špela Arhar Holdt, Ilze Auzina, Roberts Darģis, Elena Drakonaki, Jennifer-Carmen Frey, Isidora Glišić, Pinelopi Kikilintza, Lionel Nicolas, Mariana Romanyshyn, Alexandr Rosen, Alla Rozovskaya, Kristjan Suluste, Oleksiy Syvokon, Alexandros Tantos, Despoina-Ourania Touriki, Konstantinos Tsiotskas, Eleni Tsourilla, Vassilis Varsamopoulos, Katrin Wisniewski, Aleš Žagar, and Torsten Zesch. 2025b. Towards better language representation in Natural Language Processing — a multilingual dataset for text-level Grammatical Error Correction. *International Journal of Learner Corpus Research*.

Atsushi Mizumoto and Masaki Eguchi. 2023. Exploring the potential of using an AI language model for automated essay scoring. *Research Methods in Applied Linguistics*, 2(2):100050.

Tomoya Mizumoto, Yuta Hayashibe, Mamoru Komachi, Masaaki Nagata, and Yuji Matsumoto. 2012. The effect of learner corpus size in grammatical error correction of ESL writings. In *Proceedings of COLING 2012: Posters*, pages 863–872, Mumbai, India. The COLING 2012 Organizing Committee.

Ricardo Muñoz Sánchez, Simon Dobnik, Maria Irena Szawerna, Therese Lindström Tiedemann, and Elena Volodina. 2024. Did the names I used within my essay affect my score? diagnosing name biases in automated essay scoring. In *Proceedings of the Workshop on Computational Approaches to Language Data Pseudonymization (CALD-pseudo 2024)*, pages 81–91, St. Julian's, Malta. Association for Computational Linguistics.

Courtney Napoles, Keisuke Sakaguchi, Matt Post, and Joel Tetreault. 2015. Ground truth for grammatical error correction metrics. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 588–593, Beijing, China. Association for Computational Linguistics.

Courtney Napoles, Keisuke Sakaguchi, and Joel Tetreault. 2017. JFLEG: A fluency corpus and benchmark for grammatical error correction. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 2, Short Papers*, pages 229–234, Valencia, Spain. Association for Computational Linguistics.

Diane Nicholls. 2003. The Cambridge Learner Corpus – error coding and analysis for lexicography and ELT. In *Proceedings of the Corpus Linguistics 2003 conference; UCREL technical paper number 16*.

Kostiantyn Omelianchuk, Vitaliy Atrasevych, Artem Chernodub, and Oleksandr Skurzhanskyi. 2020. GECToR – grammatical error correction: Tag, not rewrite. In *Proceedings of the Fifteenth Workshop on Innovative Use of NLP for Building Educational Applications*, pages 163–170, Seattle, WA, USA → Online. Association for Computational Linguistics.

Robert Östling, Katarina Gillholm, Murathan Kurfalı, Marie Mattson, and Mats Wirén. 2024. Evaluation of really good grammatical error correction. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 6582–6593, Torino, Italia. ELRA and ICCL.

Ellis Page. 1968. The use of the computer in analyzing student essays. *International Review of Education*, 14.

Brian Riordan, Andrea Horbach, Aoife Cahill, Torsten Zesch, and Chong Min Lee. 2017. Investigating neural architectures for short answer scoring. In *Proceedings of the 12th Workshop on Innovative Use of NLP for Building Educational Applications*, pages 159–168, Copenhagen, Denmark. Association for Computational Linguistics.

Keisuke Sakaguchi, Courtney Napoles, Matt Post, and Joel Tetreault. 2016. Reassessing the goals of grammatical error correction: Fluency instead of grammaticality. *Transactions of the Association for Computational Linguistics*, 4:169–182.

Felix Stahlberg and Shankar Kumar. 2020. Seq2Edits: Sequence transduction using span-level edit operations. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 5147–5159, Online. Association for Computational Linguistics.

Ryszard Staruch. 2025. UAM-CSI at MultiGEC-2025: Parameter-efficient LLM Fine-tuning for Multilingual Grammatical Error Correction. In *Proceedings of the 14th Workshop on Natural Language Processing for Computer Assisted Language Learning*.

Toshikazu Tajiri, Mamoru Komachi, and Yuji Matsumoto. 2012. Tense and aspect error correction for ESL learners using global context. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 198–202, Jeju Island, Korea. Association for Computational Linguistics.

Elena Volodina, Yousuf Ali Mohammed, Sandra Derbring, Arild Matsson, and Beata Megyesi. 2020. Towards privacy by design in learner corpora research: A case of on-the-fly pseudonymization of Swedish learner essays. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 357–369, Barcelona, Spain (Online). International Committee on Computational Linguistics.

Elena Volodina, Christopher Bryant, Andrew Caines, Orphée De Clercq, Jennifer-Carmen Frey, Elizaveta Ershova, Alexandr Rosen, and Olga Vinogradova. 2023. MultiGED-2023 shared task at NLP4CALL: Multilingual grammatical error detection. In *Proceedings of the 12th Workshop on NLP for Computer Assisted Language Learning*, pages 1–16, Tórshavn, Faroe Islands. LiU Electronic Press.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander Rush. 2020. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45, Online. Association for Computational Linguistics.

Helen Yannakoudakis, Øistein E Andersen, Ardeshir Geranpayeh, Ted Briscoe, and Diane Nicholls. 2018. Developing an automated writing placement system for ESL learners. *Applied Measurement in Education*, 31(3):251–267.

Helen Yannakoudakis, Ted Briscoe, and Ben Medlock. 2011. A new dataset and method for automatically grading ESOL texts. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pages 180–189, Portland, Oregon, USA. Association for Computational Linguistics.

Zheng Yuan and Christopher Bryant. 2021. Document-level grammatical error correction. In *Proceedings of the 16th Workshop on Innovative Use of NLP for Building Educational Applications*, pages 75–84, Online. Association for Computational Linguistics.

A Write & Improve prompts selected for the Corpus

Table 14: The 50 Write & Improve prompts selected for inclusion in the Corpus, with their designated difficulty level, number of user-prompt essay sets and total number of essays in the Corpus.

prompt ID	level	n.essay.sets	n.essays	prompt
7fb0941f7b739572	Beginner	104	542	You receive this email from your English pen friend: Hi, It was my birthday on Saturday. I went with my family to a restaurant. After the meal they all sang Happy Birthday to me and we had a cake with candles. What do people do to celebrate birthdays in your country? Is it the same or different? Write soon! Jack Write an email to your friend.
91626718bd3e657e	Beginner	104	432	You receive this email from your English pen friend, Matt. _Hi there,_ _It is the end of the holidays and we will start school again on Monday. In this new school year, I need to work very hard because I will have exams. I must improve in Maths and History. And you? You said you want to improve your English. What will you do to improve it? Good luck!_ _Matt_ Write to Matt and tell him what you will do to improve your English in this new school year.
e4c497f7613dcf8a	Intermediate	103	513	It is the start of a new school year at your school. Your English teacher asks you to write a report about what you want to achieve in this new school year. You should write about: The areas of your learning you are pleased with. The areas of your learning you want to improve. What you will do to achieve your goals.
029cb9051e09e186	Beginner	103	497	* What do you do in your free time? How often do you do it? * What do you like about that activity and why? * Tell us about the last time you did this.
6e0152cf70e42a3a	Beginner	103	492	Your friend Leyla had a birthday party yesterday. Leyla invited you to the party but you couldn't go. Write a note to Leyla. In your note: * tell her why you couldn't go to the party * ask her what happened at the party * invite her to do something with you next weekend. Write about 70 words.
9d22a5a9df98ce09	Intermediate	103	487	Some people believe that it is good for students to have the same teacher for several years. Others think students benefit more from having new teachers each year. Which do you think is better? Give specific reasons to support your opinion.
a16c7dd5b5c59b57	Beginner	103	483	* Do you like to exercise? How often do you exercise? * When you exercise, do you prefer exercising at a gym or outside? Why? * Tell us about the kinds of activities you like to do for exercise.
6c812937c853f95a	Intermediate	103	471	Your friend Alex has invited you to a party this Saturday afternoon but you can't go. Write an email to Alex. In your email: * apologise to Alex * explain why you can't go * suggest another day when you could meet. Write your **email** in about **100 words.**
79d7b6720d30aa96	Beginner	103	465	Write an email to your English friend. In your email, tell your friend about your favourite film. What kind of film is it? What is it about? Do you think your friend would like the film?

…continued

prompt ID	level	n.essay.sets	n.essays	prompt
394b6f53ecc6f98d	Intermediate	103	461	Video games are very popular with children. However, some parents feel that video games can have a negative impact on their children, but others believe that they may have some positive effects. How do you feel about children playing video games? Give specific reasons and examples to support your opinion.
e0d347fb2afea90b	Advanced	103	447	Some university students move away and live on their own while they are attending school, but others live at home with their family while they are students. Which do you think is a better choice? Give reasons to explain your answer.
9564cdfd1ab47024	Advanced	103	447	Most medical doctors want to work in cities where they will make a lot of money. Yet in many rural areas there are not enough doctors. What are the reasons for this problem? How can it be solved?
e3b16d9c3c7b3cc8	Beginner	103	420	Read this email from your cousin, Ana. Hi, I saw your mum yesterday. She told me you are learning English. That's good news! But I want to know why you are studying English. Why is learning the English language important to you? I wish you luck! Ana **Write an email to Ana and answer her question.**
6d735aae1b87fea9	Intermediate	102	548	You have had a class discussion on the importance of exercise. Now your teacher has asked you to write a composition answering the following question: _Exercise is important for our health, but are there things that are more important?_ Write your composition.
30d627bea13ce859	Advanced	102	537	Would you like to live a long, calm life until 100 years old or would you prefer a shorter life but much more intense. Decide and explain your reasons.
6748f85e7e8f36b6	Intermediate	102	506	Some students choose to study in another country for a time during high school or university. What do you think are the advantages and disadvantages of studying abroad?
29081c0eb3c823b4	Intermediate	102	504	Some people believe zoos are good places for people to learn about animals. Other people believe animals belong in nature and it is wrong to keep them in zoos. What do you think? Explain, giving specific reasons for your choice.
87f7c87fb31b5da2	Beginner	102	504	Write to a student in another country describing a special food you like to eat. Make sure to tell them: * what the food is * when you usually eat the food * why you like this food Include more details. Write at least one paragraph.
39db840b1dee1ac1	Advanced	102	496	In some offices, employees work in enclosed workstations or cubicles. Other offices have a more open design, in which there are no walls separating employees from their coworkers. Which type of arrangement do you think is most effective? Explain, giving reasons to support your answer.
23ed369963fb84a1	Beginner	102	491	Read and answer the question below. Write at least 2 paragraphs. Some people sleep in a dark and quiet room. Other people sleep with a small light on and don't mind noise. How do you like to sleep? Explain your answer.

. . . continued

prompt ID	level	n.essay.sets	n.essays	prompt
a48191f7b89e3523	Intermediate	102	480	You have had a class discussion on the purpose of television. Now your teacher has asked you to write a composition answering the following question: _Television entertains, but can it also educate?_ Write your composition.
5be9aa25d9f61171	Advanced	102	468	In some schools, students must study music and learn to play a musical instrument. In other schools, studying music may not be required or there may be no music classes at all. Do you think all students should have to study music? Support your opinion with reasons and examples.
8e0f8cbbba514658	Advanced	102	461	Some high schools have a requirement that students must complete a certain amount of volunteer work in order to graduate. What do you think about this type of requirement? Give details to support your opinion.
2865cfff040b541d	Advanced	102	445	Your town council is thinking about banning the use of mobile phones on public transport and in shops and restaurants. Write a letter to the town council giving your opinion about this and why you think it is or is not a good idea.
17eab621f9e89332	Beginner	102	432	* How often do you use a computer? * What do you usually use a computer for? * Tell us about what you enjoy doing most with a computer.
6305e66e9a1bd3db	Beginner	102	420	Last week you went on a trip with some students from your class. Write an email to your friend, Billal. In your email: * say where you went * say what you did * say why you enjoyed the trip. Write about 80 words.
98e191fa5bded725	Beginner	102	411	A friend from your class, Cristina, had an accident last week. Now, her doctor says she must stay at home for two weeks. Write a note to Cristina. In the note: * tell her you are sorry about her problem * say what you can do to help her * say when you can visit her. Write about 50 words.
4671630d1fa03592	Beginner	102	410	Your friend Rie sends you an email: _Hi, Can you come to my house for dinner tomorrow? What do you like to eat? Come at 7 o'clock. See you. Rie_ Write a note to Rie. In your note: * say you can come to dinner * say what you like to eat. Write about 30 words.
58b9f5f522c0a1ba	Intermediate	102	404	You have been asked to write a short **article** for your school magazine with the title: _My reasons for learning English_ Describe your reasons for learning English **and** explain why it is important to you. **Write your article.**
2adc00d65002bc78	Beginner	101	566	* What is your favorite indoor activity? * Why do you like this activity? * Tell us about the last time that you did this activity.
1a893b704d9d98d4	Advanced	101	546	Many people believe that it is important to always tell the truth. In your opinion, is it ever acceptable to tell a lie? Explain your answer giving reasons and specific examples.
49f744b86a7de069	Beginner	101	483	Your neighbour is going to look after your house while you are on holiday. Write a note to tell him what he must do. For example, you can write about: * keys * plants * pets * mail

. . . continued

prompt ID	level	n.essay.sets	n.essays	prompt
43448f28d3ce551a	Beginner	101	474	Your English friend Chris is coming with you and your family to the local museum next Saturday. Write an email to Chris. Tell Chris: - how you will travel to the museum - what he can see at the museum - what he needs to bring. Write 25 words or more.
9975b4bd24cfacd8	Intermediate	101	420	Write about a holiday that you went on that was really important to you. Say where you went and what you saw, and explain why it was the best holiday you ever had.
13020c2a757e39d5	Intermediate	101	405	You have recently had a class discussion about the importance of taking care of your health. Now your English teacher has asked you to write an essay giving your opinions on the following statement: _It is more important to be healthy than to be rich._ Write your **essay.**
a3271cf10d32656b	Intermediate	101	384	Is learning to speak a foreign language as important these days as it was in the past? Write an **essay** giving your opinion.
d5dbd9c04a5c932d	Intermediate	100	449	You have seen this notice in an international music magazine: **Articles wanted** **Do we need music in our lives?** Are there times when music is important? How does music make us feel? What would life be like without music? Write and tell us what you think. Write your **article** in **140 – 190** words in an appropriate style.
58852905890136d4	Intermediate	100	429	You see a request for articles in a magazine: We invite you to write about: _The Future of Travel_. In what ways do you think travel will be different in the future? In what ways might it still be the same? The writer of the best article will win a prize! Write your article.
91752a41d038711e	Beginner	100	403	Your friend Sam is coming to your house tomorrow evening. Write an email to Sam. Tell Sam: * **what** time to come * **what** to bring * **how** to get to your house. Write **25** **words** or more.
40037060180c75db	Beginner	100	402	Your teacher asks you to write about what you do in your free time. Write about: * the things you do in your free time * when you do these things. Write about 40 words.
de206ba1795e8dc5	Intermediate	99	503	What is your favourite film? Write a letter to your pen friend. Explain why you like the film and why you think your pen friend should watch it.
046678dfe42e3068	Intermediate	99	500	Your English teacher told you about Write and Improve and said you should use it to practise and improve your English writing. Your teacher has asked you to write **a report** about your experience with Write and Improve, describing what you liked or didn't like about it and why. Finish by saying whether you would recommend your classmates to try it. Write your **report**.
305b51299790d8a8	Intermediate	99	461	You have been asked to write a short __article__ for your school magazine with the title: *My favourite season of the year.* Describe your favourite season __and__ explain why it is so special for you. __Write your article.__

…continued

prompt ID	level	n.essay.sets	n.essays	prompt
179938ed5e90b50a	Intermediate	99	455	You see a request for articles in a magazine: We invite you to write about: _The School of the Future_. In what ways do you think schools will be different in the future? In what ways might they still be the same? The writer of the best article will win a prize! Write your article.
719435a48902a4b5	Beginner	98	390	You are studying English. You also want to study another subject. Write a note to your teacher. In your note: * say what other subject you want to study * say why you want to do this course * say when you want to do this course * ask your teacher to help you.
2ef4e350c1009968	Beginner	97	451	Write a paragraph for a brochure about your home town. You must include TWO of these things: · Information about **important landmarks** · Information about **shopping** and **accommodation** · Information about **parks, museums** and **restaurants** · Information about **night-time entertainment.**
6cb0c85d60a4ef2e	Advanced	96	559	Because of the negative effects on health of an inactive lifestyle, all university students and school children should be forced to do sports activities at least 3 times per week. Do you agree? Write your **argumentative essay**.
82f61f70f00cce70	Intermediate	96	392	Should computers and tablets completely replace books in schools and universities? Write an **essay** giving your opinion.
e404501e9bbbc094	Advanced	95	470	A well-known magazine for English language learners has asked for reviews from readers about useful resources for learning English. You have recently discovered Write and Improve, a new writing tool for learners of English. Write a review for the magazine describing the main features of the tool and why it is/isn't helpful. Include details about what you like about the tool and about any difficulties you had with it, and finish by giving your recommendation for other readers. Write your **review**.
d3e35405e2378ac9	Intermediate	92	400	Last week, you went to a music concert. Now you are writing a letter to an English-speaking friend. Explain why you decided to go to the concert, give some information about who was performing, and describe what happened at the concert.

B Corpus metadata

The metadata for every essay in the Corpus, contained in the file `en-writeandimprove2024-corpus.tsv`, features the following columns:

- **public_essay_id** – essay identifier;
- **created_epoch** – time of essay creation in epoch seconds;
- **created_timestamp** – date and time of essay creation, GMT;
- **public_prompt_id** – prompt identifier;
- **public_user_id** – user identifier;
- **user_prompt** – user and prompt identifiers concatenated;
- **essay_version_num** – version number within a user-prompt essay set;
- **is_first_version** – is the first version within a user-prompt essay set, T/F;
- **is_final_version** – is the final version within a user-prompt essay set, T/F;
- **language** – first language of the user;
- **text** – the essay;
- **wi_suspecttokens** – character indices of any automatically detected errors in a comma-separated list of start+end pairs (otherwise NA);
- **automarker_cefr_level** – CEFR level estimated by the W&I auto-marker;
- **humannotator_cefr_level** – CEFR level assigned by human experts, final versions only (otherwise NA);
- **split** – train/dev/test split assigned to the first and final versions only (otherwise NA).